

Computer Networks

@CS.NYTU

Lecture 2: Applications

Instructor: Kate Ching-Ju Lin (林靖茹)

Slides modified from

“Computer Networking: A Top-Down Approach” 7th Edition

Outline

- Principles of network applications
- Web and HTTP
- E-Mail and SMTP
- DNS
- Peer-to-peer applications
- **Video streaming and CDN**

“Multimedia” Networking

text

CS, NCTU
Multimedia
Networking
2017 Spring
ker ker

image



video




gaming



VR




audio / music



Why Video Streaming Important?

- **High demand**

- Over 50% Internet traffic



- **High bandwidth requirement**


- 30, 60, 120 frames per seconds
- Video rate: 100kbps – 3Mbps or even $>10\text{Mbps}$ (4K)

- **Complex compression algorithms**

- **No need to receive all the bits**
- Video quality is related to packet loss rate
 - But not linear proportional

1. Video-On-Demand Streaming

- Request **on demand**
- Online downloading and playing



MOD
中華電信

You **Tube**


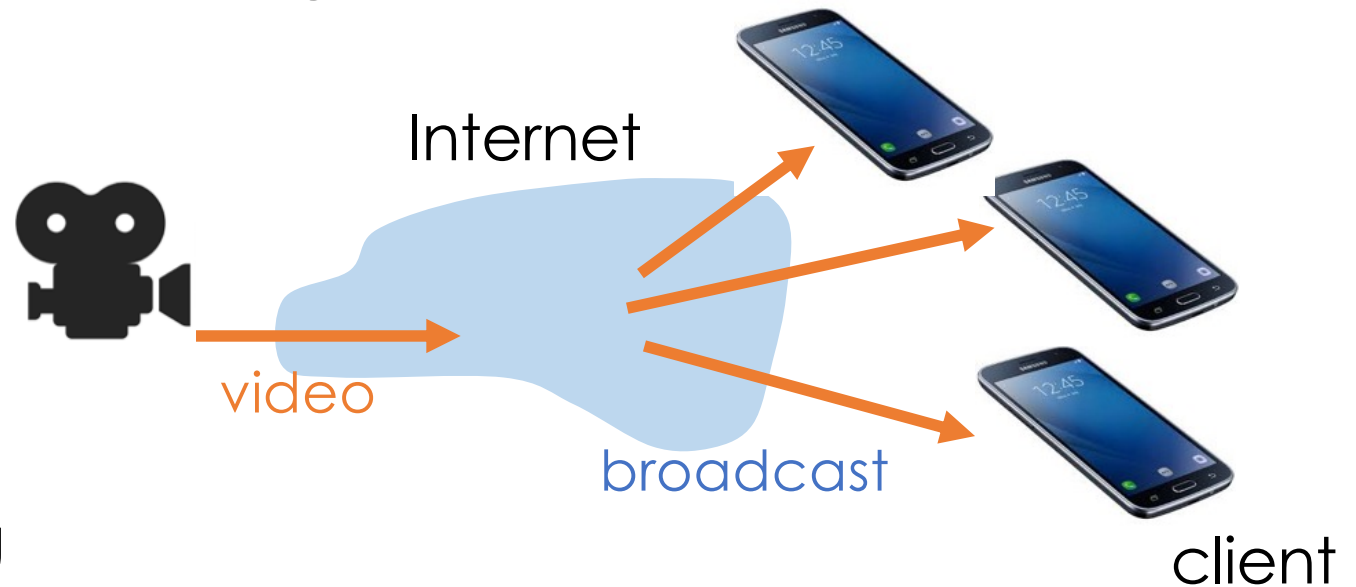
NETFLIX

2. Real-Time Video Streaming

- Record and deliver live video immediately
- Usually broadcasting



video
recording



3. HTTP Streaming

- **DASH**: *D*ynamic, *A*daptive *S*teaming over *H*TT*P*
- **server**:
 - Divide video file into multiple chunks
 - Each chunk stored, encoded at different rates
 - **Manifest file**: provides URLs for different chunks/rates
- **client**:
 - Periodically measures server-to-client bandwidth
 - Consulting manifest, requests one chunk at a time
 - Choose maximum coding rate sustainable given current bandwidth
 - Can choose different coding rates at different points in time (depending on available bandwidth at time)

Why HTTP Streaming

- Traditional streaming delivered over RTP/UDP
- However, in today's Internet, content objects are stored **Content Delivery Networks (CDN)**
 - Many CDNs do not support RTP streaming
 - RTP often does not allow traffic through firewall
 - Each RTP stream is a separate session → not scalable
- Benefits of HTTP streaming
 - Firewall friendly
 - No need to maintain the session state on the server
 - Has been standardized as "ISO/IEC 23009-1, also known as MPEG-DASH" in Apr. 2012

Why CDNs?

- **Challenge:** how to stream content (selected from millions of videos) to hundreds of thousands of *simultaneous* users?
 - **Option 1:** single, large “mega-server”
 - Reliability: single point of failure
 - Not enough bandwidth: point of network congestion
 - Far from users: long path to distant clients
 - multiple copies of video sent over outgoing link
- quite simply: this solution *doesn't scale*

Why CDNs?


- **Challenge:** how to stream content (selected from millions of videos) to hundreds of thousands of *simultaneous* users?
- **Option 2:** store/serve multiple copies of videos at geographically distributed sites (CDN)
 - **enter deep:** push CDN servers deep into many access networks
 - ✓ close to users
 - ✓ used by Akamai, 1700 locations
 - **bring home:** smaller number (10's) of larger clusters in POPs near (but not within) access networks
 - ✓ used by Limelight

Content Distribution Networks

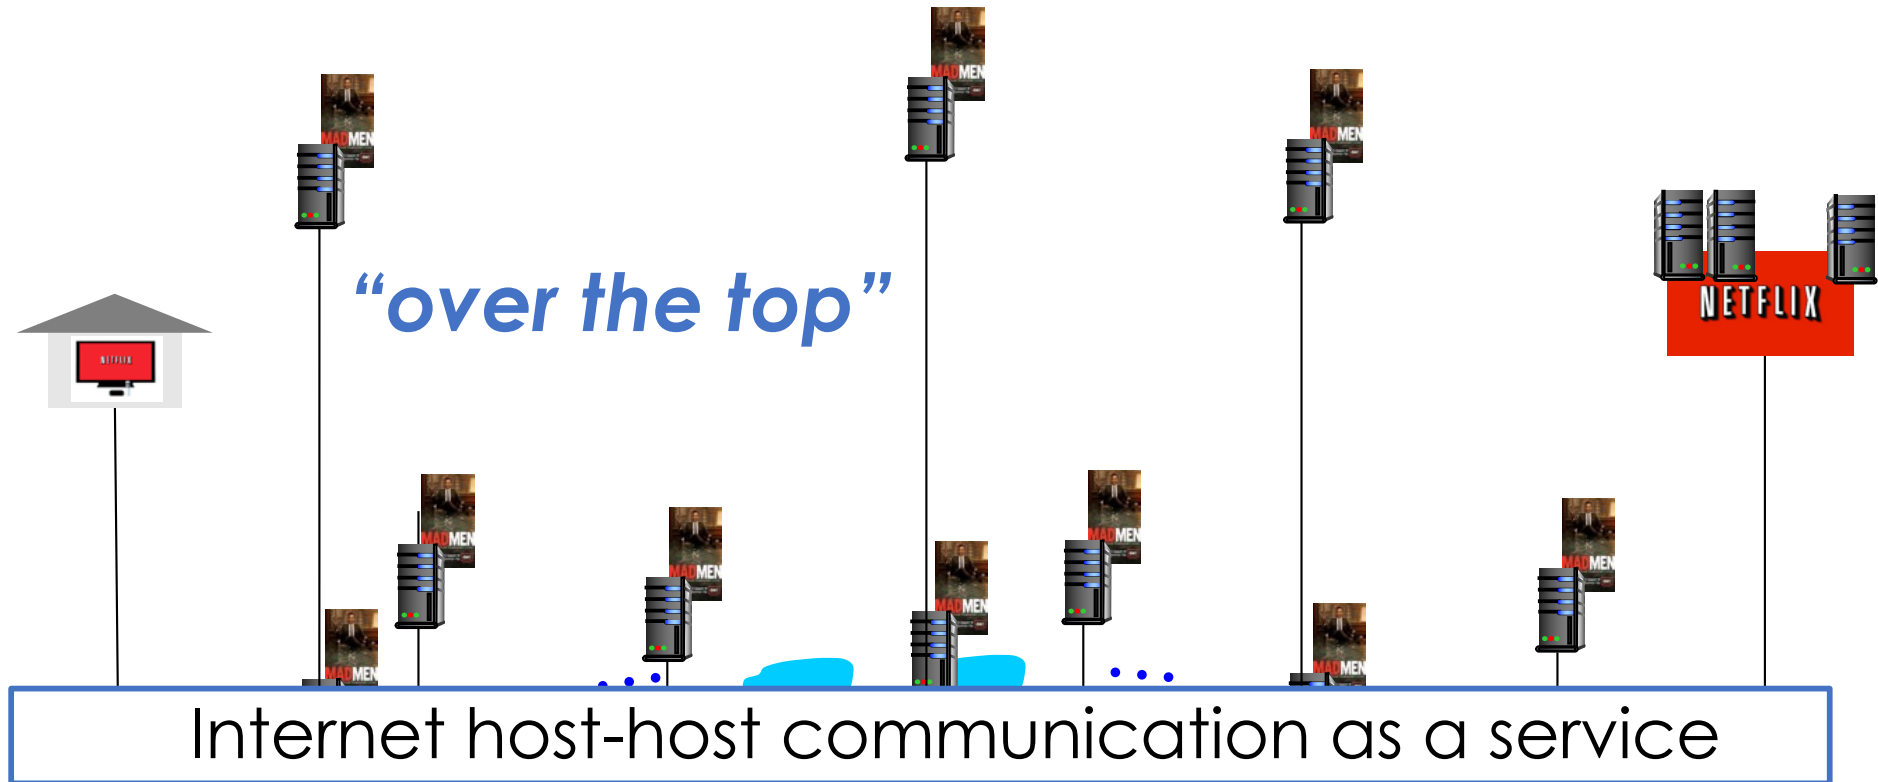
- **CDN:** an application overlay (e.g., Akamai)
- Design Space
 - **Caching** (data-driven, passive)
 - explicit
 - transparent (hijacking connections)
 - **Replication** (pro-active)
 - server farms
 - geographically dispersed (CDN)
- Three Main CDN Providers (in North America, Europe):
 - Akamai, Limelight, Level 3 CDN

Key Idea of CDN

- CDN: stores copies of content at CDN nodes
 - e.g. Netflix stores copies of *MadMen*
- Subscriber requests content from CDN
 - directed to nearby copy, retrieves content
 - may choose different copy if network path congested



Framework of CDN

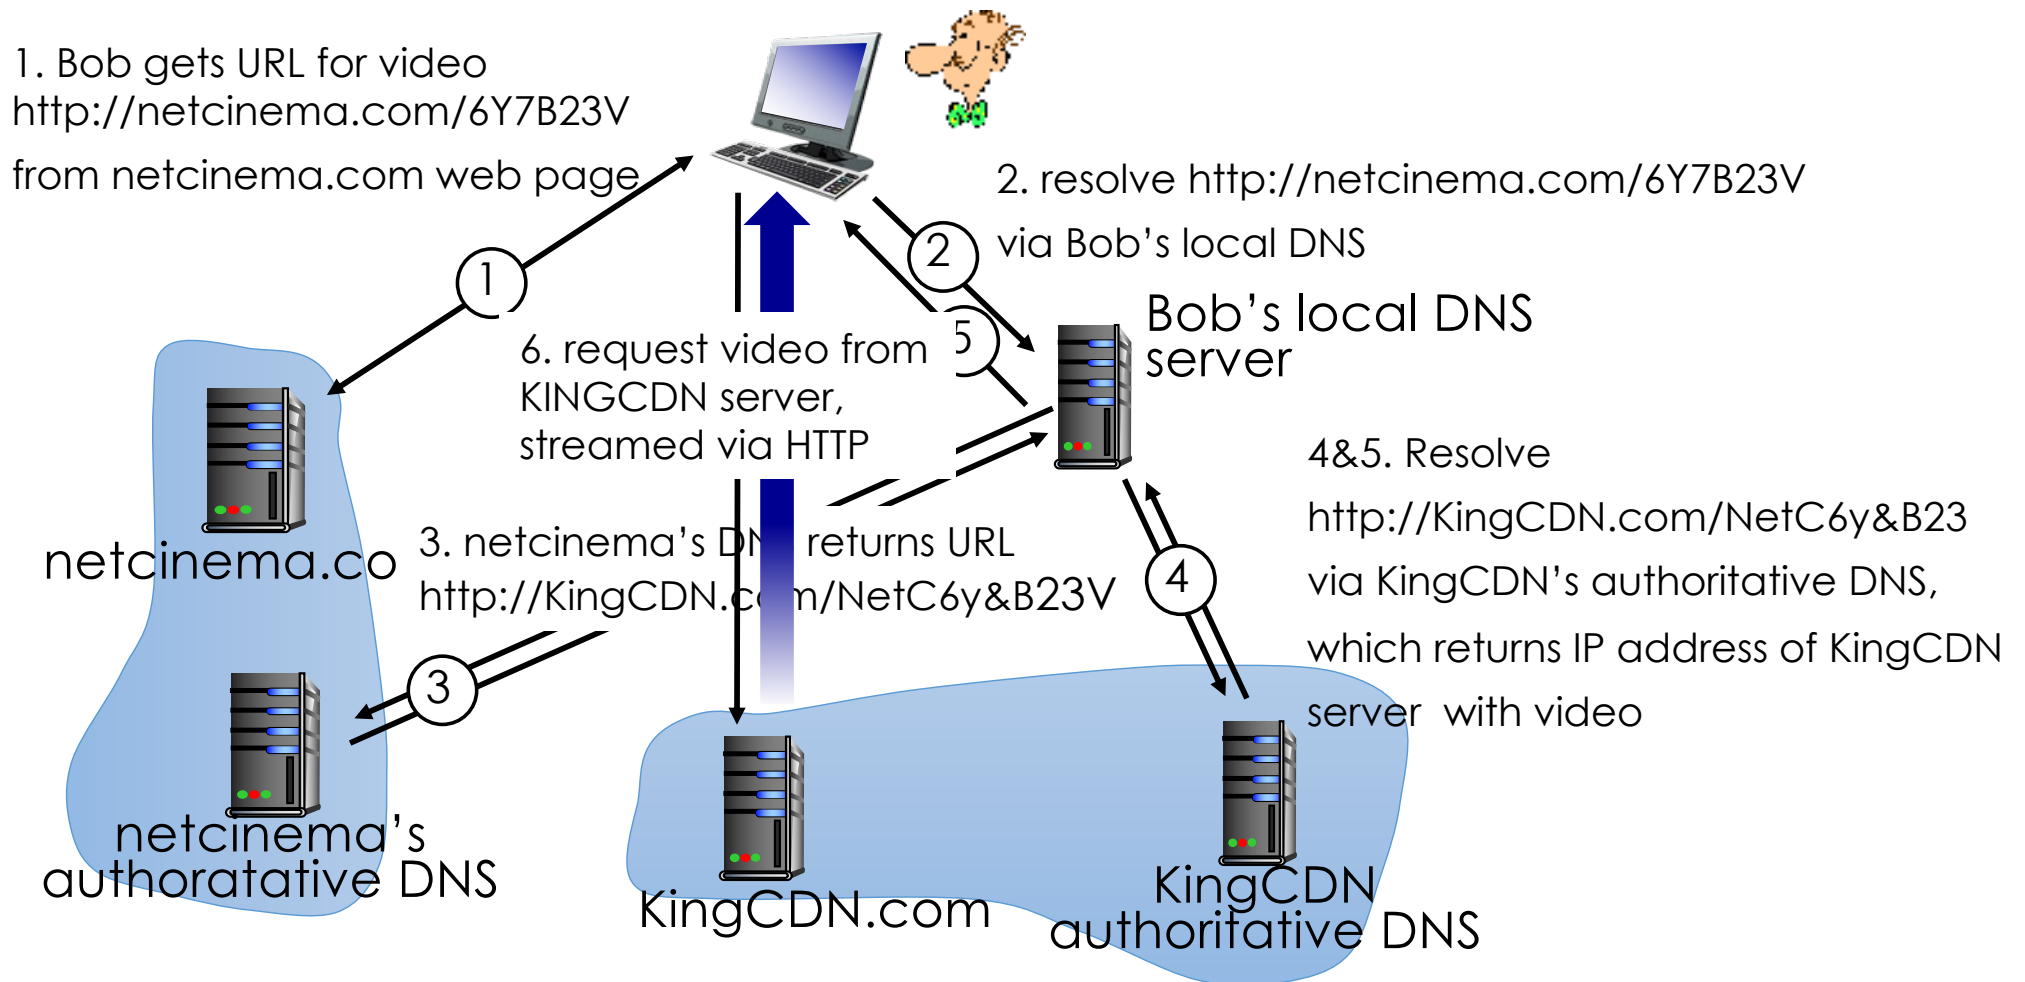


OTT challenges: coping with a congested Internet


- from which CDN node to retrieve content?
- viewer behavior in presence of congestion?
- what content to place in which CDN node?

CDN Content Access

- Bob (client) requests video <http://netcinema.com/6Y7B23V>
- video stored in CDN at <http://KingCDN.com/NetC6y&B23V>



Case Study: Netflix



Cluster Selection Strategies

- **Geographically closest**

- Geographical distance \neq path length
- Local DNS location \neq user location
- Do not explicitly consider bandwidth and delay (varying with the congestion level)

- **Fastest response time**


- Periodically probe from CDN clusters to DNS servers
- Ingest huge probing traffic
- Some DNS servers configured to “not reply”

Outline

- Principles of network applications
- Web and HTTP
- E-Mail and SMTP
- DNS
- **Peer-to-peer applications**
- Video streaming and CDN


Pure P2P Architecture

- No always-on server
 - Might have some **super nodes**, mainly used for management
- Arbitrary end systems (peers) directly communicate
 - **Self-scalable**
- Peers are intermittently connected and change IP addresses dynamically
- Examples:
 - File distribution (BitTorrent)
 - Streaming (PPstream)
 - VoIP (Skype)



P2P: Challenges

- How to locate your peer & find what you want?
- Need some kind of “directory” or “look-up” service



- centralized
- distributed, using a hierarchal structure
- distributed, using a flat structure
- distributed, without structure (“flooding”)
- distributed, using “hybrid” structured/unstructured

P2P: Challenges

- **Technical**

- **Scale**: up to hundred of thousands or millions of machines
- **Dynamics**: machines can come and go any time

- **Social, economic and legal**

- **Incentive** Issues: free-loader problem
- Vast majority of users are **free-riders**
- Most share no files and answer no queries
- A few individuals contributing to the “public good”
- **Copyrighted** content and privacy
- Trust & **security** issues

BitTorrent: Popular P2P App

- Designed for large file (e.g., video) distribution
- Focused on *efficient fetching*, not search
 - Distribute same file to many peers
 - Single publisher, many downloaders
- Divide large file into many pieces (chunks)
 - Replicate different pieces on different peers
 - A peer with a complete piece can trade with others
 - Peer can (hopefully) assemble the entire file
- Allows simultaneous downloading
 - Retrieve different pieces from different peers simultaneously
- Usually need to prevent “free loading”

BitTorrent Components

- **Seed**

- Peer with entire file
- Fragmented in pieces

- **Leacher**

- Peer with an incomplete copy of the file

- **Torrent file**

- Passive component
- Store summaries of the pieces to allow peers to verify their integrity

- **Tracker**

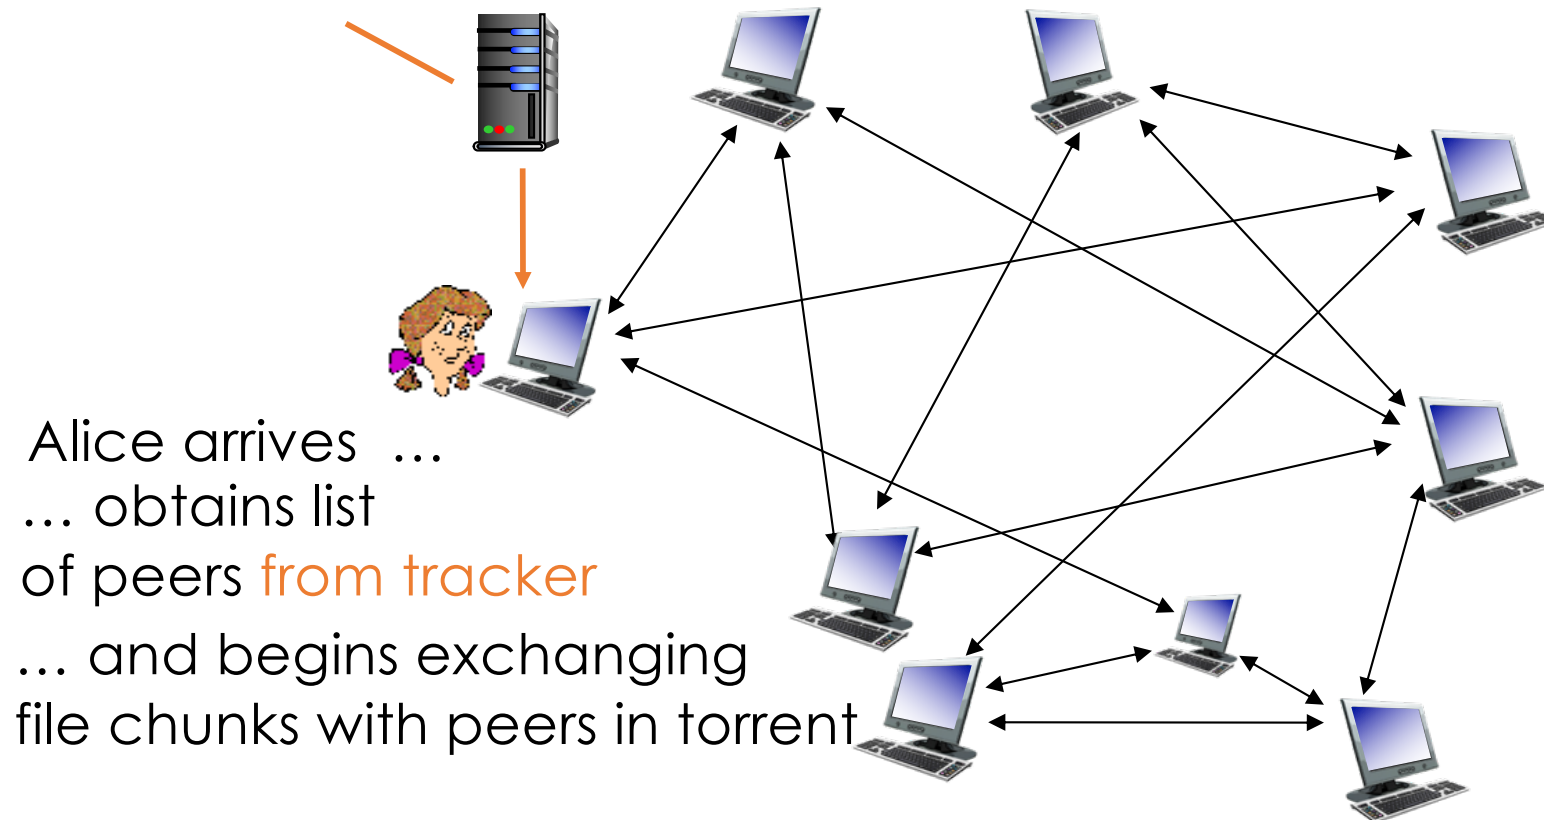
- Allows peers to find each other
- Returns a list of random peers

BitTorrent Protocol

- File divided into 256Kb chunks
- Peers in torrent send/receive file chunks

tracker: tracks peers participating in torrent

torrent: group of peers exchanging chunks of a file




BitTorrent Protocol

- Peer joining torrent:
 - Has no chunks, but will accumulate them over time from other peers
 - Register with tracker to get a list of peers (“neighbors”)
 - Exchange chunks with neighbors
- Peer may change neighbors over time
 - Tracker periodically suggests some random peers
 - Keep those random peers as neighbors if they are better (in terms of bandwidth or content)
- **Churn**: peers may come and go
- Once peer has entire file, it may (selfishly) leave or (altruistically) remain in torrent

BitTorrent: Tit-for-Tat

- (1) Alice “optimistically unchokes” Bob
- (2) Alice becomes one of Bob’s top-four providers; Bob reciprocates
- (3) Bob becomes one of Alice’s top-four providers



BitTorrent: Scheduling

Two choices should be made

- **Which chunk to download?**

- **rarest first**: download the chunk with the least copies first
- Why? More precious and might become disappear in the network
- Tracker should provide such information

- **Where to download?**

- **Highest rate first**: from a neighbor who can provide a higher downloading rate
- Update the neighbor list if the random peer suggested by the tracker supports a higher rate

BT: requesting, sending file chunks

requesting chunks:

- at any given time, different peers have different subsets of file chunks
- periodically, Alice asks each peer for list of chunks that they have
- Alice requests missing chunks from peers, rarest first

sending chunks: tit-for-tat

- Alice sends chunks to those peers currently sending her chunks at highest rate
 - other peers do not receive chunks from Alice
 - re-evaluate top 4 every 10 secs
- every 30 secs: randomly select another peer, starts sending chunks
 - “optimistically unchoke” this peer
 - newly chosen peer may join top 4